



COSC 3360/6310 FIRST ASSIGNMENT

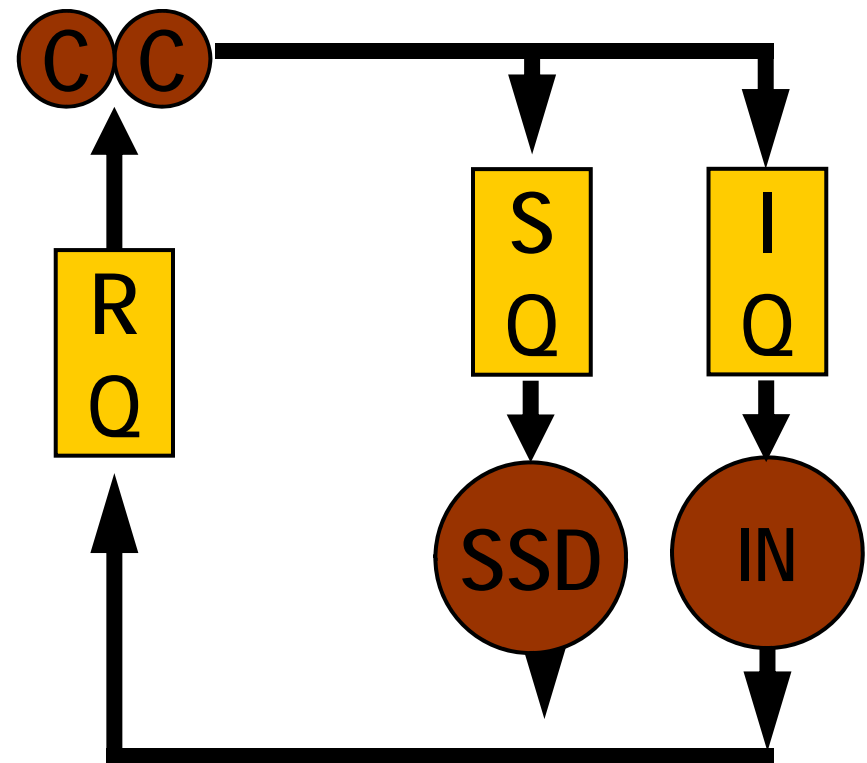
Spring 2018



The model

We have

- One multi-core CPU
- One SSD
- One input device
- Three queues
 - CPU queue
"ready queue"
 - SSD queue
 - Input queue





FIRST EXAMPLE



Start: Process 0 at $t = 5\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

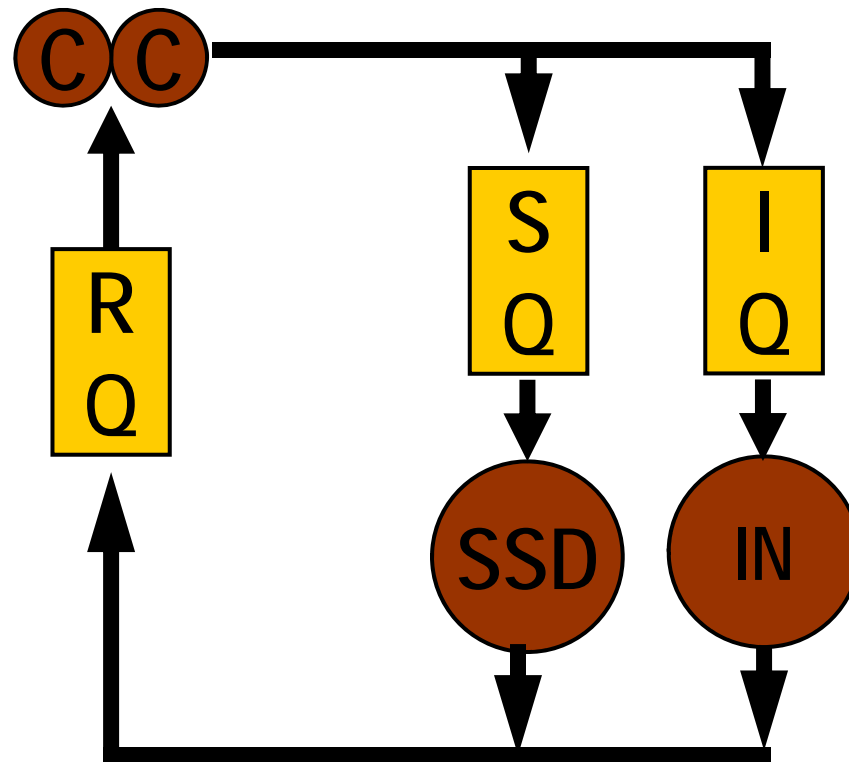
CORE 30

NEW 100

CORE 20

SSD 0

CORE 20





Your program will display

- Process 0 starts
Time = 5 ms

or

- Process 0 starts at $t = 5$ ms



Process 0 gets a core at $t = 5\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

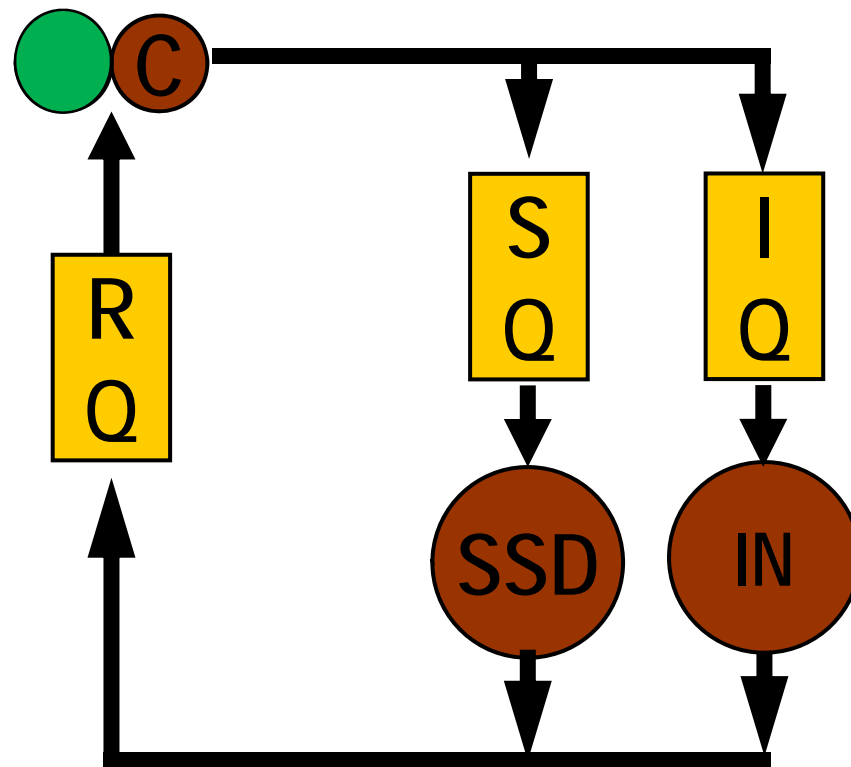
NEW 100

CORE 20

SSD 0

CORE 20

First core busy until $t = 105\text{ms}$



What happens next?

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

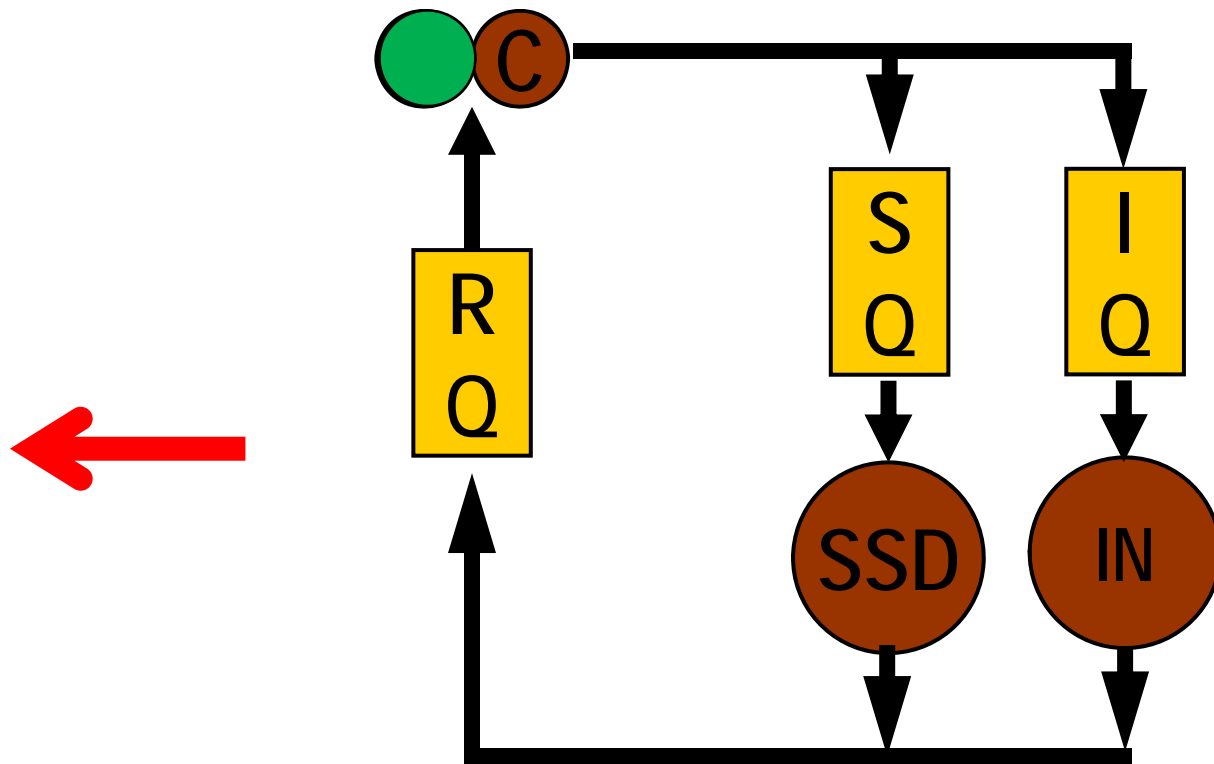
NEW 100

CORE 20

SSD 0

CORE 20

First core busy until $t = 105$ ms





Process 1 starts at $t = 100\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

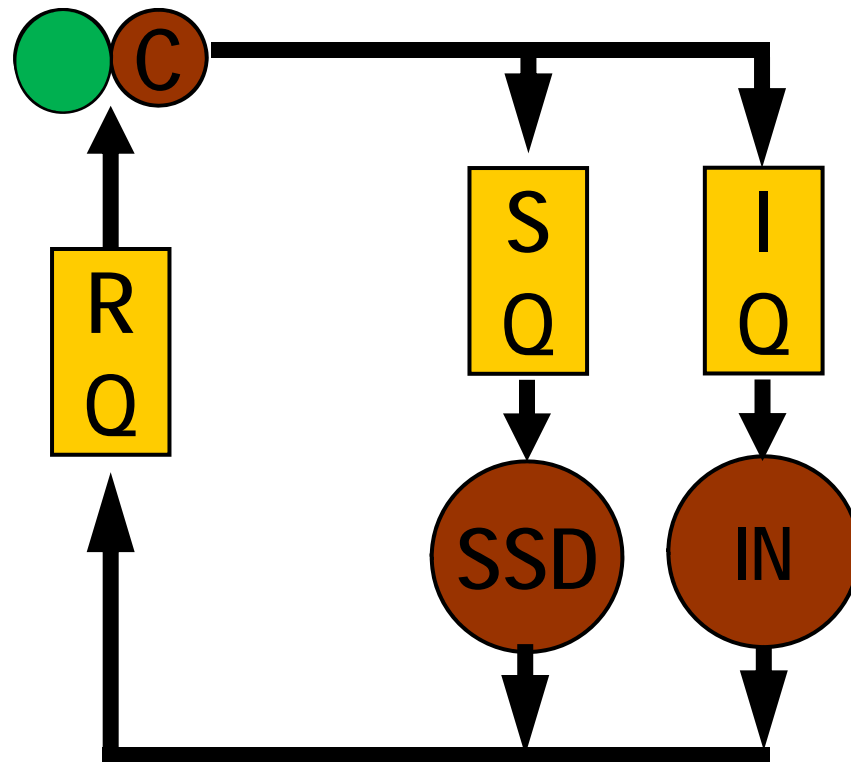
NEW 100

CORE 20

SSD 0

CORE 20

First core busy until $t = 105\text{ms}$





Your program will display

- Process 1 starts
Time = 100 ms
Process 0 is RUNNING

or

- Process 1 starts at $t = 100$ ms
Process 0 is RUNNING

Process 1 gets core at $t = 100\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100

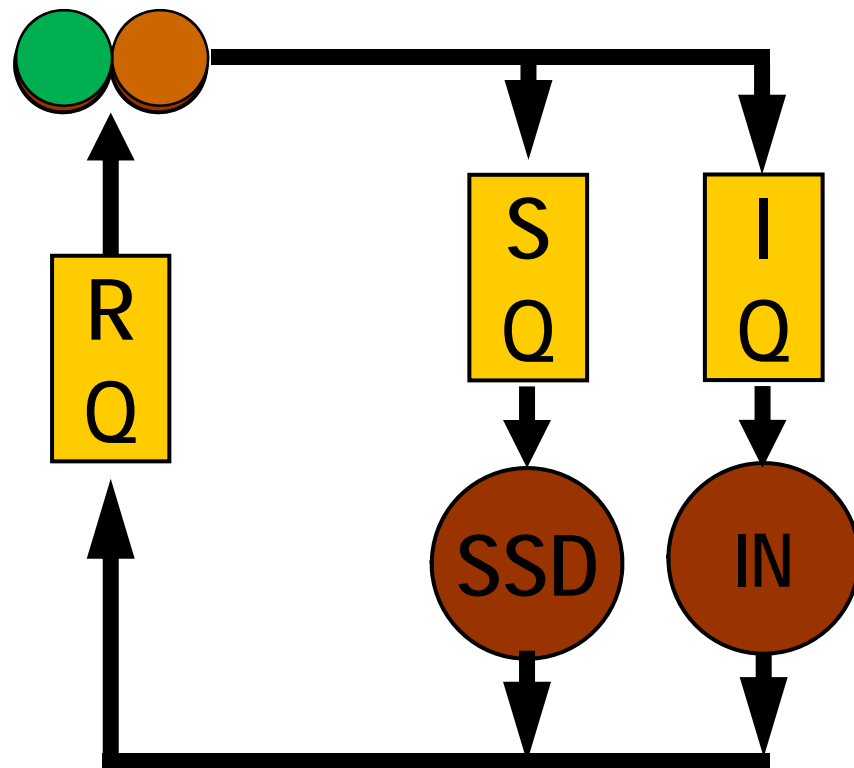
CORE 20

SSD 0

CORE 20

First core busy until $t = 105\text{ ms}$

Second core busy until $t = 120\text{ms}$

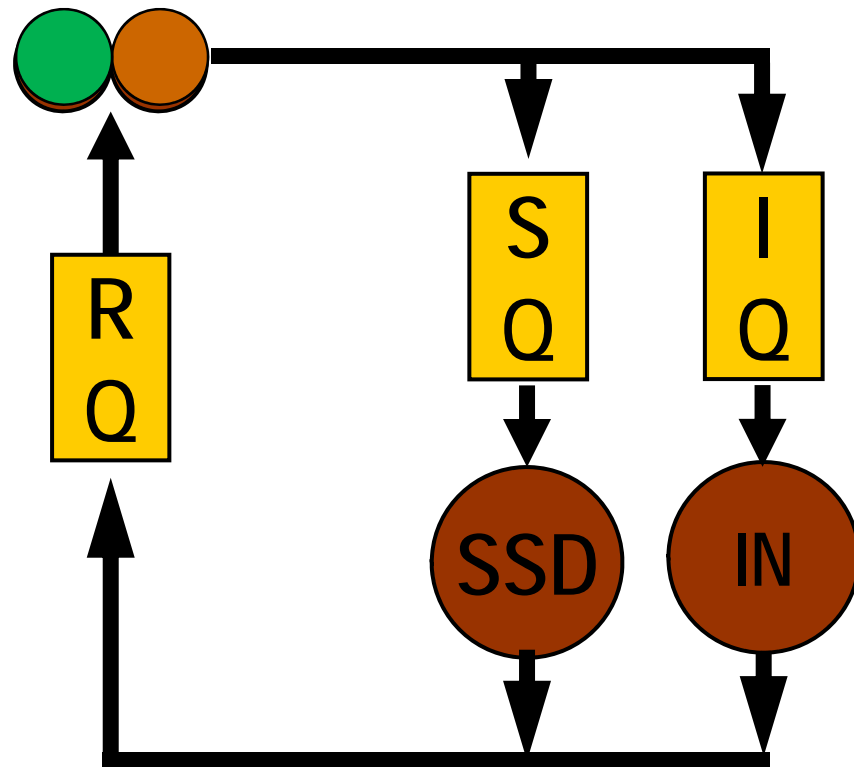




What happens next?

NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20

First core busy until $t = 105$ ms
Second core busy until $t = 120$ ms



Process 0 starts I/O at $t = 105\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100

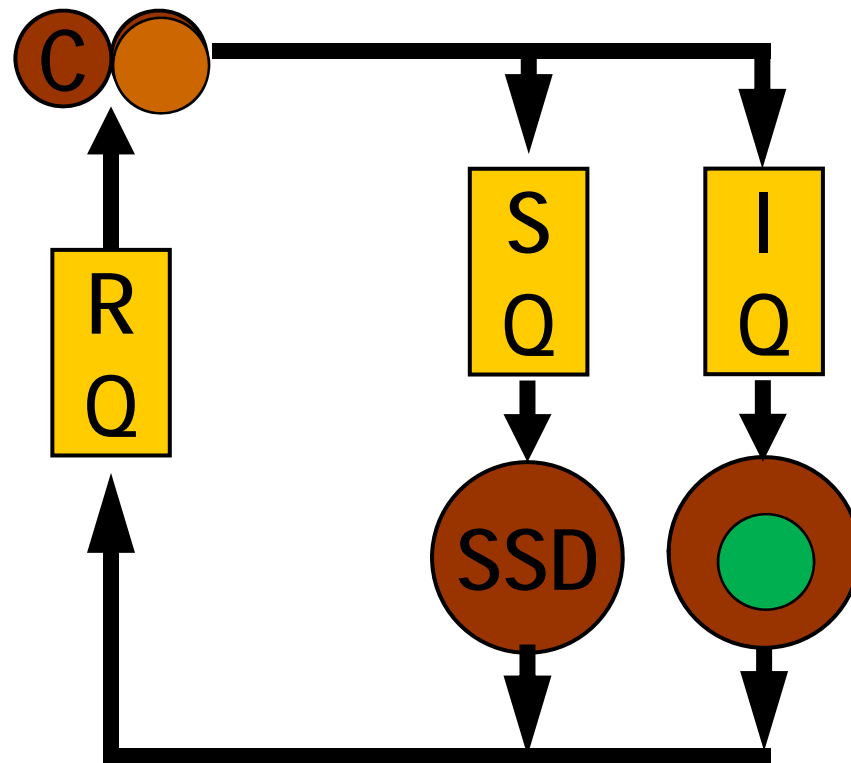
CORE 20

SSD 0

CORE 20

Second core busy until $t = 120\text{ms}$

User busy until $t = 5105\text{ms}$



What happens next?

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100


CORE 20

SSD 0

CORE 20

Second core busy until $t = 120\text{ms}$

User busy until $t = 5105\text{ms}$





Process 1 gets SSD at t = 120ms

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100


CORE 20

SSD 0

CORE 20

SSD busy until t = 120 ms


User busy until t = 5105 ms



What is next?

NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20

SSD busy until $t = 120$ ms ←
User busy until $t = 5105$ ms





Process 1 gets core at $t = 120\text{ms}$

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100

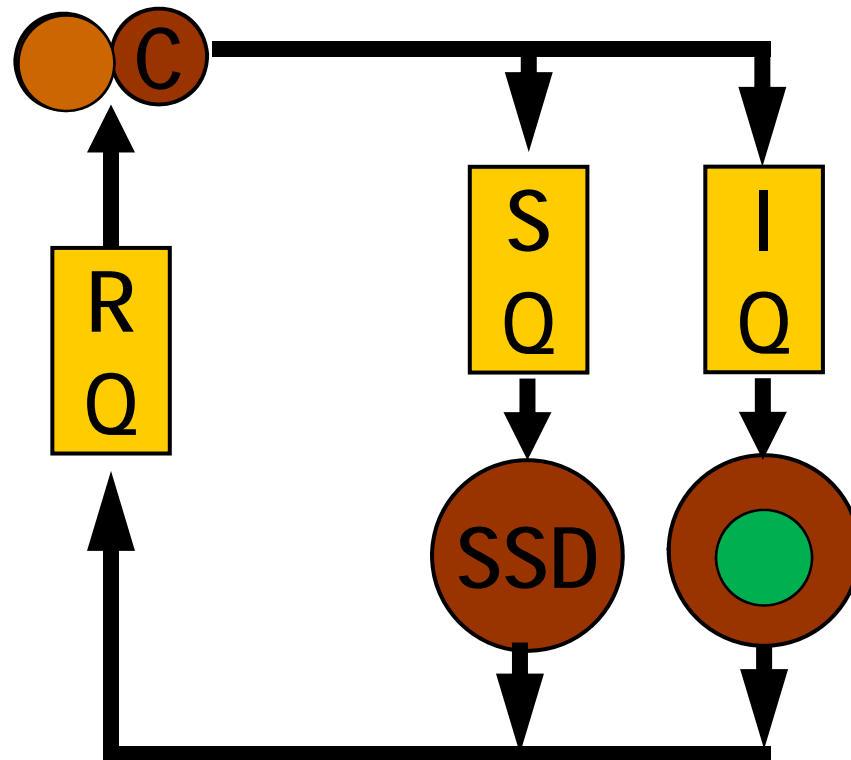
CORE 20

SSD 0

CORE 20

First core busy until $t = 140\text{ms}$

User busy until $t = 5105\text{ms}$



What happens next?

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100

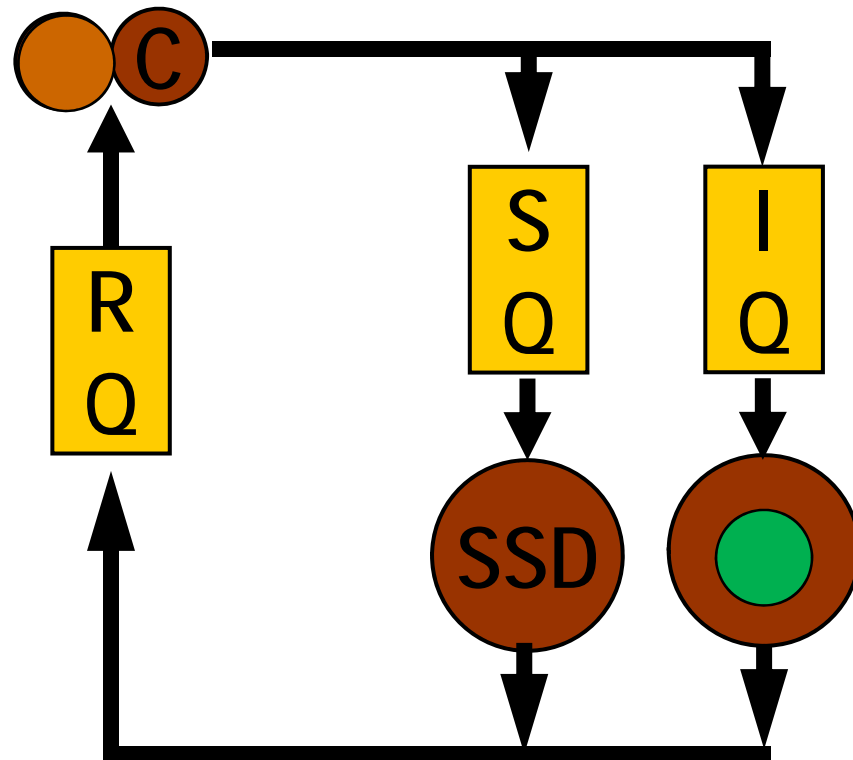
CORE 20

SSD 0

CORE 20

First core busy until $t = 140$ ms ←

Process 0 busy until $t = 5105$ ms



What happens next?

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

NEW 100

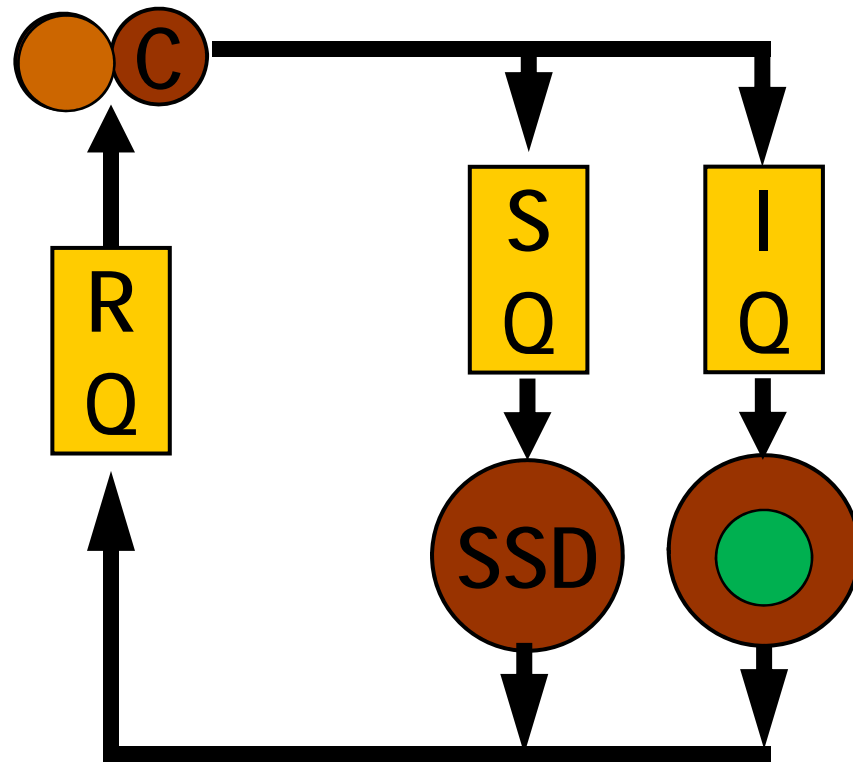
CORE 20

SSD 0

CORE 20

First core busy until $t = 140$ ms ←

Process 0 busy until $t = 5105$ ms



Process 1 ends at t = 140ms

NCORES 2

User busy until t = 5105 ms

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

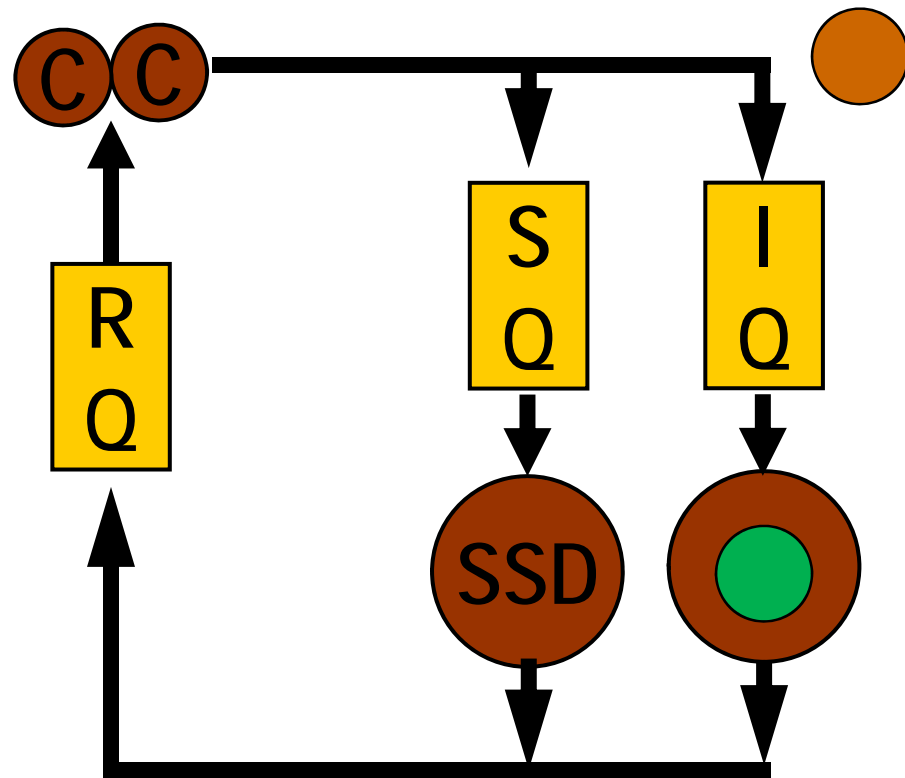
CORE 30

NEW 100

CORE 20

SSD 0

CORE 20





Your program will display

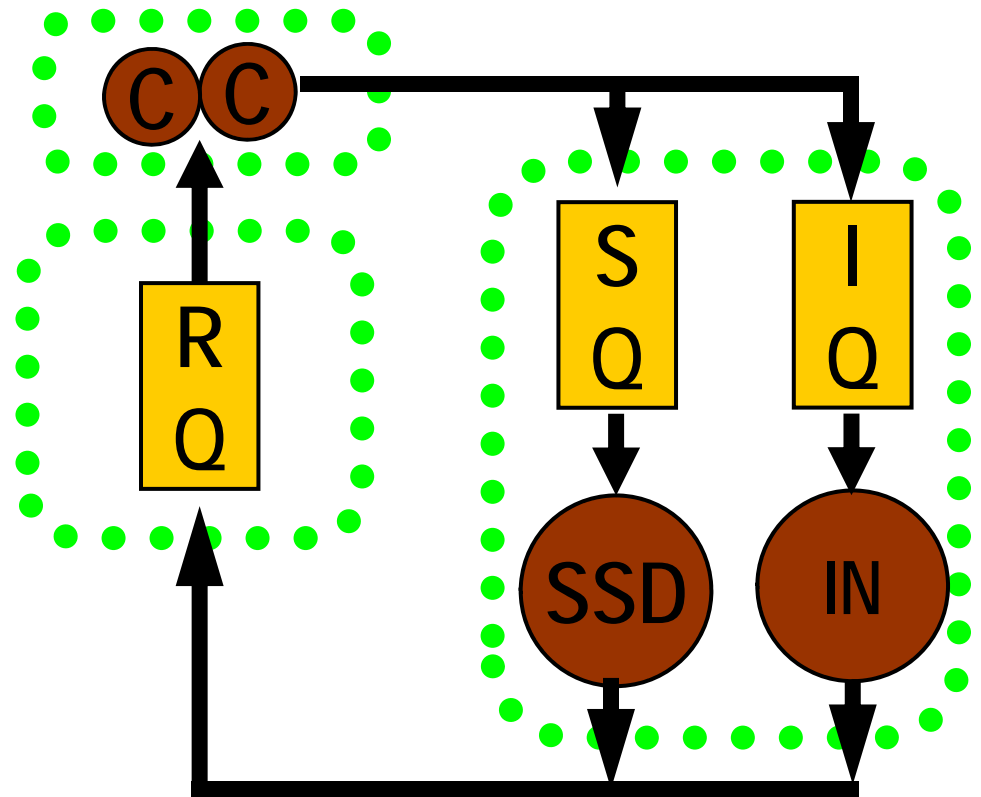
- Time = 140 ms
Process 1 terminates
Process 0 is BLOCKED

or

- Process 1 terminates at $t = 140$ ms
Process 0 is BLOCKED

The three process states

- **RUNNING** means executing on a core
- **READY** means waiting for a core
- **BLOCKED** means being neither **RUNNING** or **READY**



What happens next?

NCORES 2

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

CORE 30

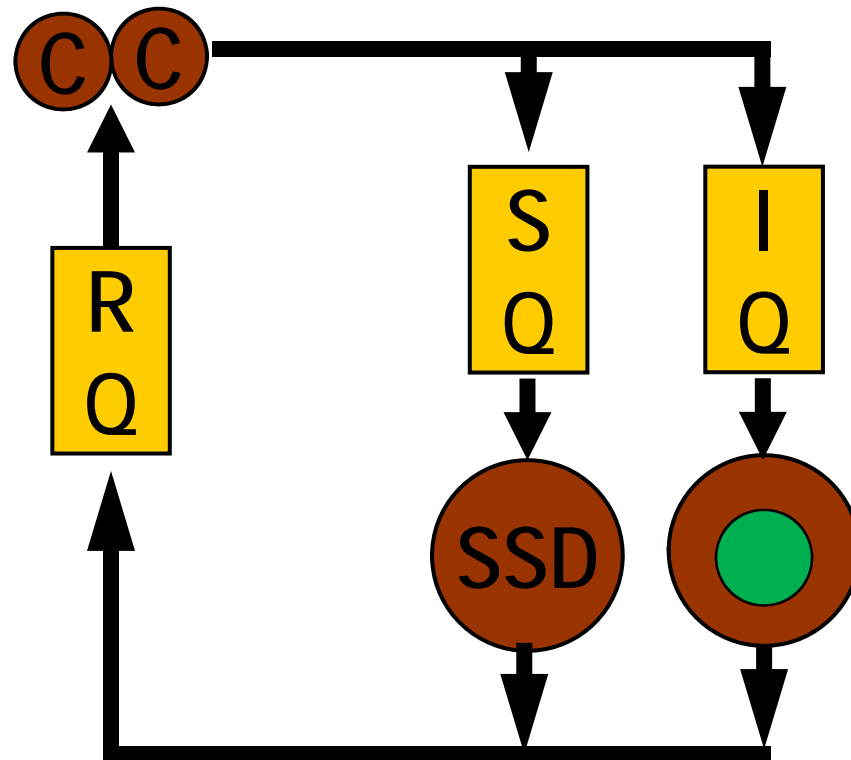
NEW 100

CORE 20

SSD 0

CORE 20

User is busy until $t = 5105$ ms ←



Process 0 gets core at t=5105ms

NCORES 2

Core busy until t = 5185 ms

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

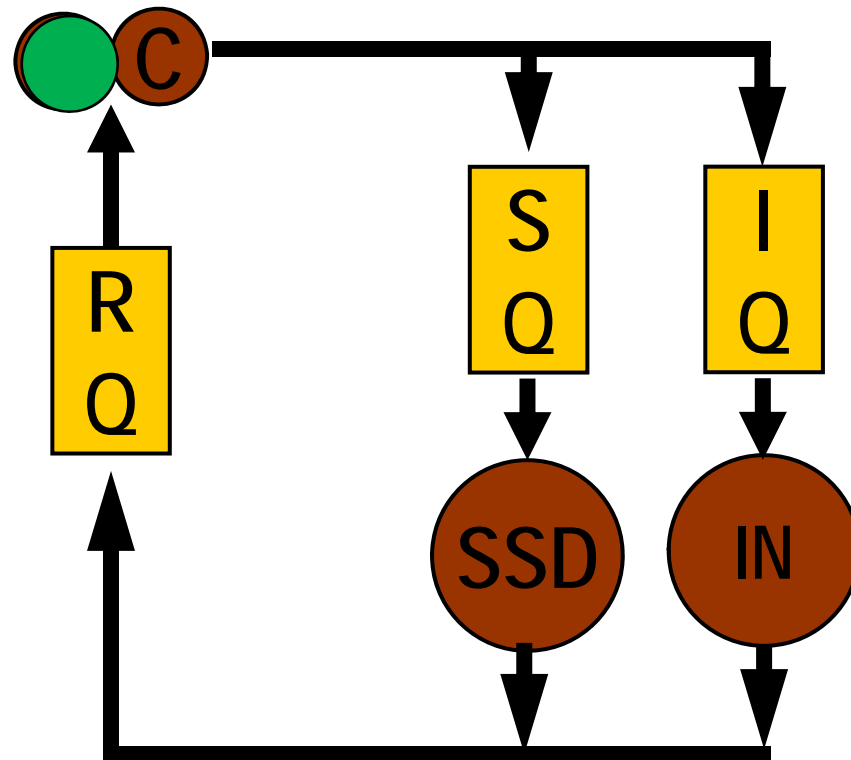
CORE 30

NEW 100

CORE 20

SSD 0


CORE 20



What happens next?

NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20

Core busy until $t = 5185$ ms ←



Process 0 gets SSD at t=5185ms

NCORES 2

SSD is busy until t = 5186 ms

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1


CORE 30

NEW 100

CORE 20

SSD 0


CORE 20



What happens next?

NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20

SSD busy until $t = 5186$ ms ←



Process 0 gets core at t=5186ms

NCORES 2

A core is busy until t = 5216ms

NEW 5

CORE 100

INPUT 5000

CORE 80

SSD 1

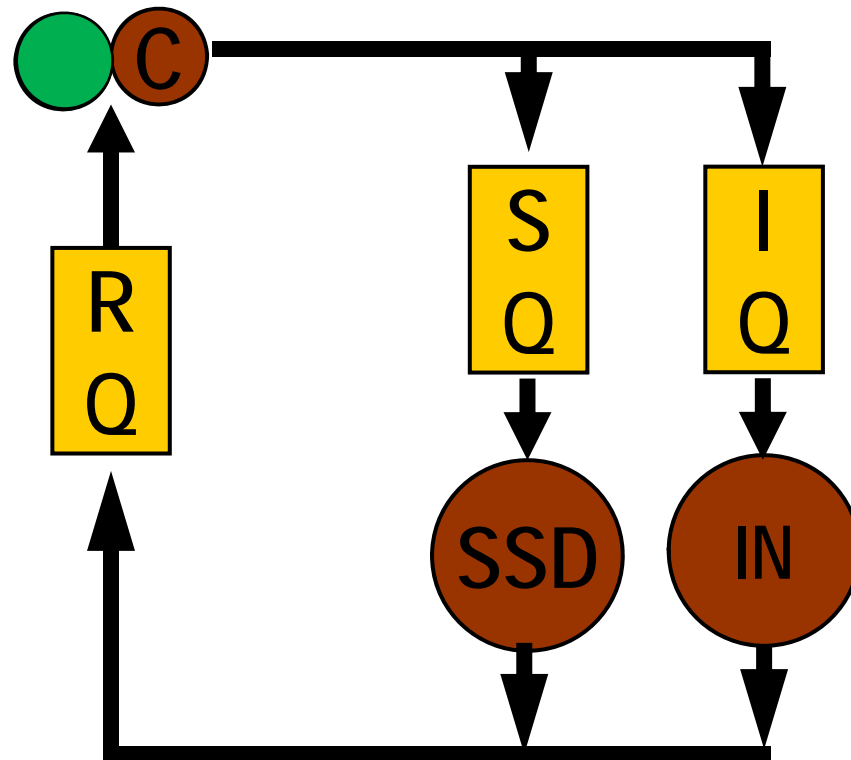
CORE 30

NEW 100

CORE 20

SSD 0

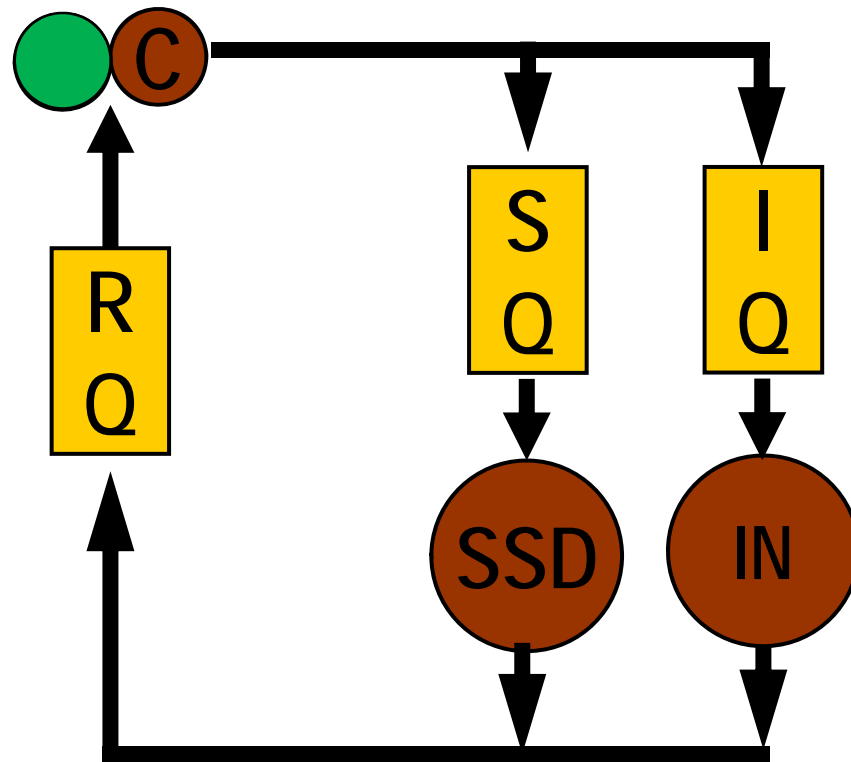
CORE 20



What happens next?

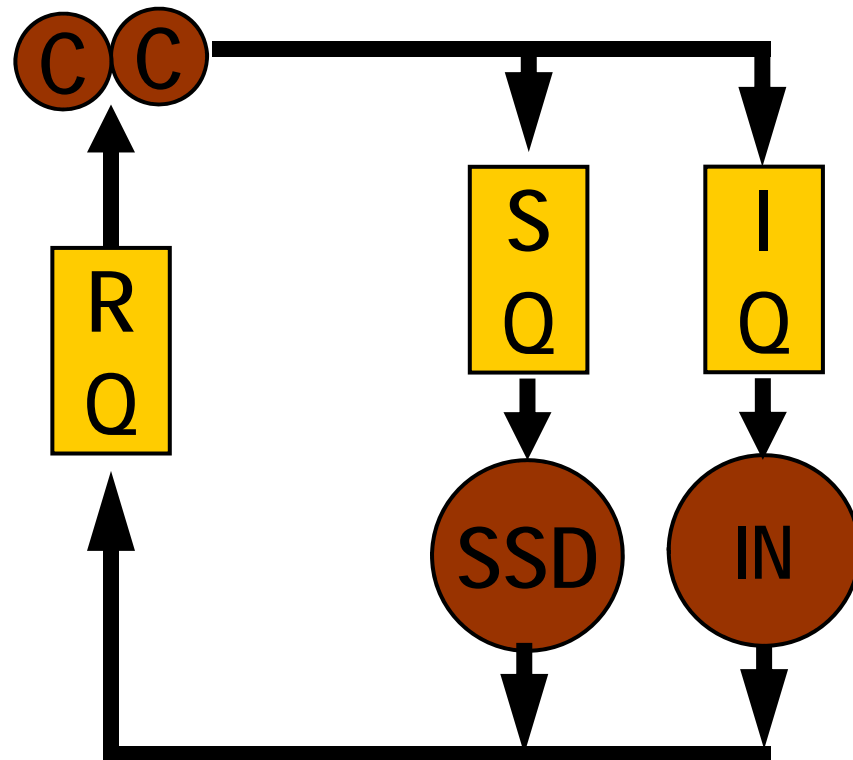
NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20

A core is busy until $t = 5216\text{ms}$ ←



Process 0 terminates

NCORES 2
NEW 5
CORE 100
INPUT 5000
CORE 80
SSD 1
CORE 30
NEW 100
CORE 20
SSD 0
CORE 20





Your program will display (I)

- Process 0 terminates
Time = 5216 ms

or

- Process 0 terminates at $t = 5216$ ms



Your program will display (II)

■ SUMMARY:

Number of processes that completed: 2

Total number of SSD accesses: 2

Average SSD access time: 0.5 ms

Total elapsed time: 5216 ms

Core utilization: 4.79 percent

SSD utilization: 0.02 percent



How to compute core utilization

- Keep track of total time for all RUN requests:
 - $100 + 80 + 30 + 20 + 20 = 250$ ms
- Divide by elapsed time:
 - $250/5216 = \mathbf{0.0479}$ (rounded)
- Since NCORES = 2 , the maximum CORE utilization is 2.0



How to compute SSD utilization

- Keep track of total time for all SS requests:
 - $0 + 1 = 1$ ms
- Divide by elapsed time:
 - $1/5216 = \mathbf{0.0002}$ (rounded)
- SSD utilization is normally computed for each storage device, so it will never exceed 1.0



INTRODUCING CONTENTION

P0 gets a core at $t = 5\text{ms}$

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50

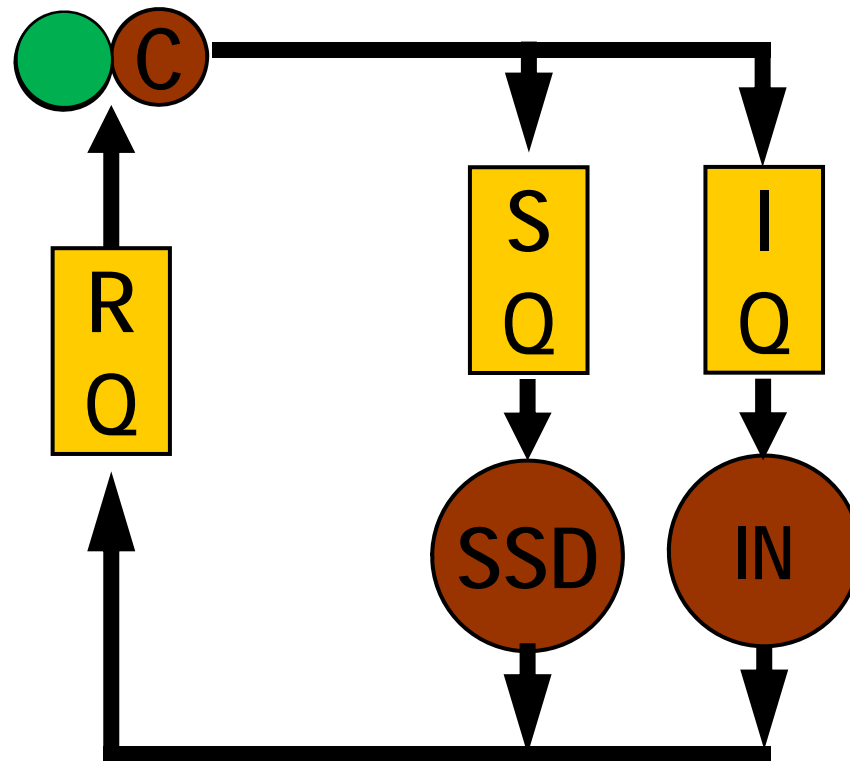
NEW 30

CORE 20

SSD 0

CORE 20

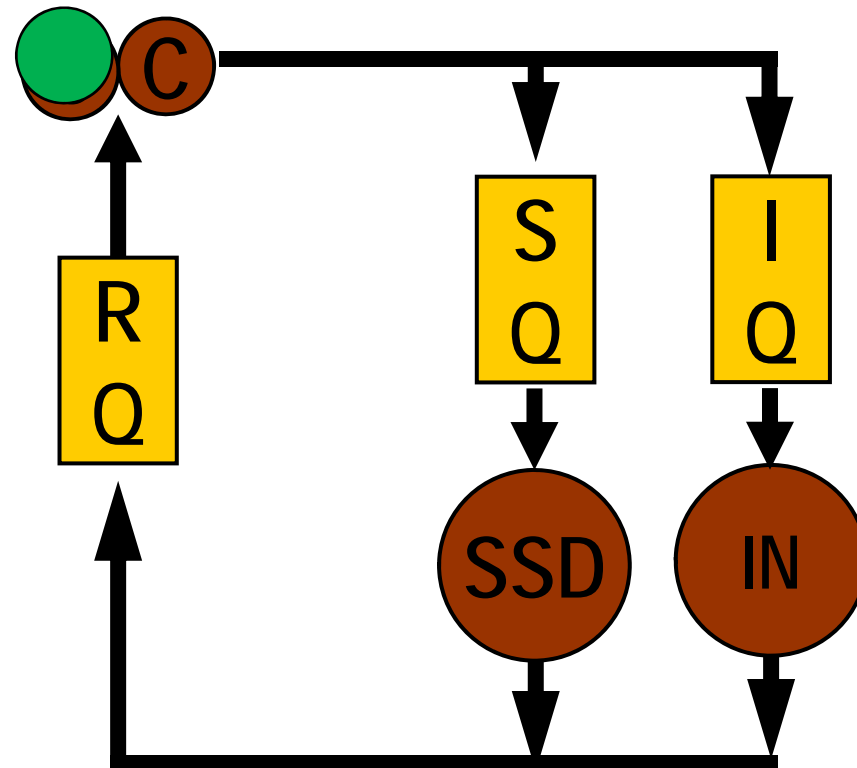
P0 holds a core until $t = 105\text{ms}$



What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20 ←
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

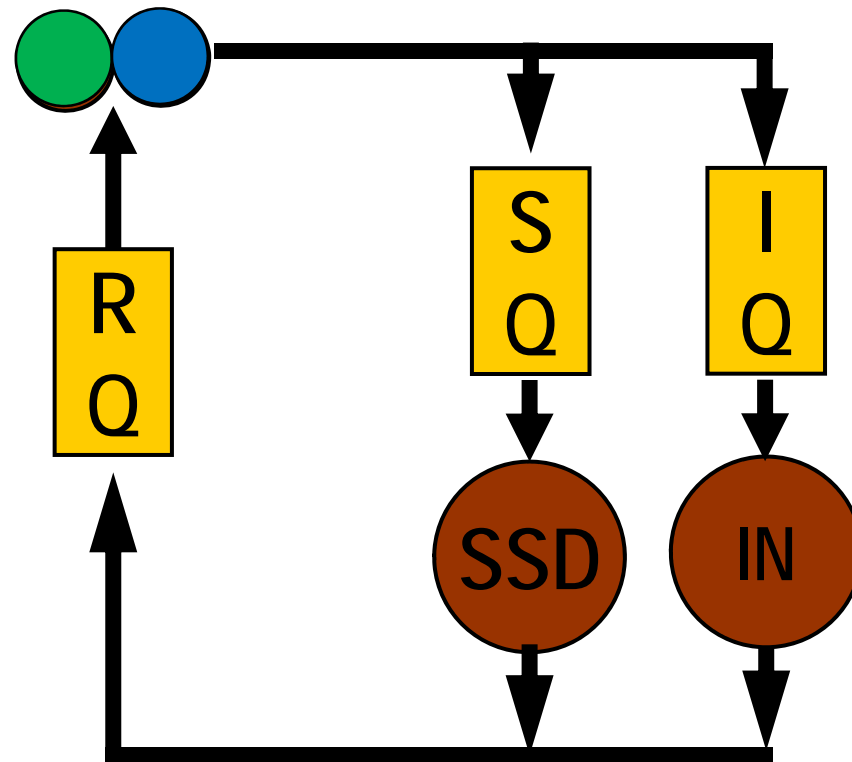
P0 holds a core until $t = 105$ ms



P1 gets a core at $t = 20\text{ms}$

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

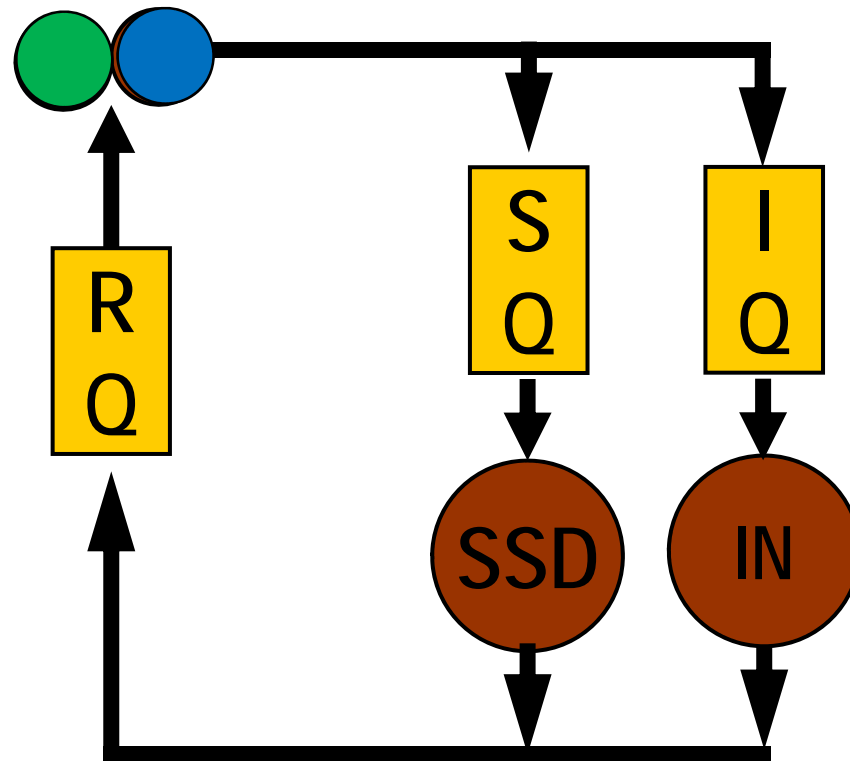
P0 holds a core until $t = 105\text{ms}$
P1 holds a core until $t = 70\text{ms}$



What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30 ←
CORE 20
SSD 0
CORE 20


P0 holds a core until $t = 105\text{ms}$
P1 holds a core until $t = 70\text{ms}$



P2 waits for a core

```
NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20
```

First core busy until $t = 105\text{ms}$
Second core busy until $t = 70\text{ms}$




What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 105\text{ms}$

P1 holds a core until $t = 70\text{ms}$



P1 gets SSD at $t = 70\text{ms}$
P2 get core

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50

NEW 30

CORE 20


SSD 0

CORE 20

P0 holds a core until $t = 105\text{ms}$

P2 holds a core until $t = 90\text{ms}$

P1 holds SSD until $t = 70\text{ms}$




What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 105\text{ms}$

P2 holds a core until $t = 90\text{ms}$


P1 holds SSD until $t = 70\text{ms}$ ←



Process 1 waits for a core

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20


P0 holds a core until $t = 105\text{ms}$
P2 holds a core until $t = 90\text{ms}$



What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20


P0 holds a core until $t = 105\text{ms}$
P2 holds a core until $t = 90\text{ms}$ ←



P2 gets SSD at time $t = 90\text{ms}$
P1 gets a core

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 105\text{ms}$
P1 holds a core until $t = 140\text{ms}$
P2 holds SSD until $t = 90\text{ms}$




What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 105\text{ms}$

P1 holds a core until $t = 140\text{ms}$


P2 holds SSD until $t = 90\text{ms}$ ←



P2 waits for a core

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20


P0 holds a core until $t = 105\text{ms}$
P1 holds a core until $t = 140\text{ms}$



What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 105\text{ms}$
P1 holds a core until $t = 140\text{ms}$



P0 gets SSD at $t = 105\text{ms}$
P2 gets a core

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50

NEW 30

CORE 20


SSD 0

CORE 20

P1 holds a core until $t = 140\text{ms}$

P2 holds a core until $t = 125\text{ms}$

P0 holds SSD until $t = 105\text{ms}$




What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P1 holds a core until $t = 140\text{ms}$

P2 holds a core until $t = 125\text{ms}$


P0 holds SSD until $t = 105\text{ms}$



P0 waits for a core

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20


P1 holds a core until $t = 140\text{ms}$
P2 holds a core until $t = 125\text{ms}$



What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P1 holds a core until $t = 140\text{ms}$
P2 holds a core until $t = 125\text{ms}$ ←



P2 terminates

P0 gets a core at time $t = 125\text{ms}$

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50

NEW 30

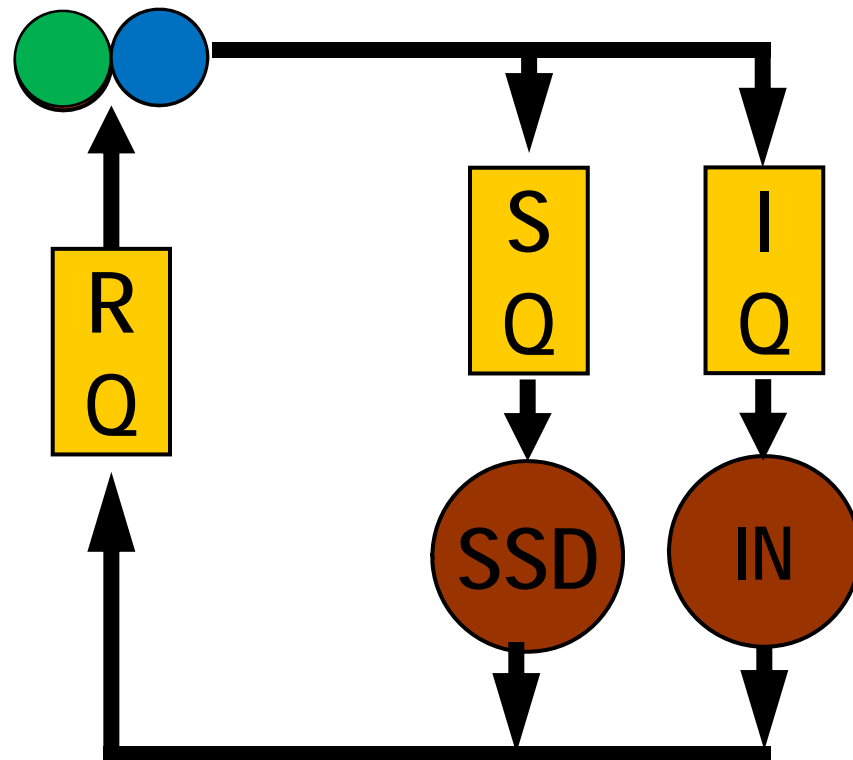
CORE 20

SSD 0

CORE 20

P1 holds a core until $t = 140\text{ms}$

P0 holds a core until $t = 155\text{ms}$





Your program will display

- Time = 125 ms
Process 2 terminates
Process 0 is RUNNING
Process 1 is RUNNING

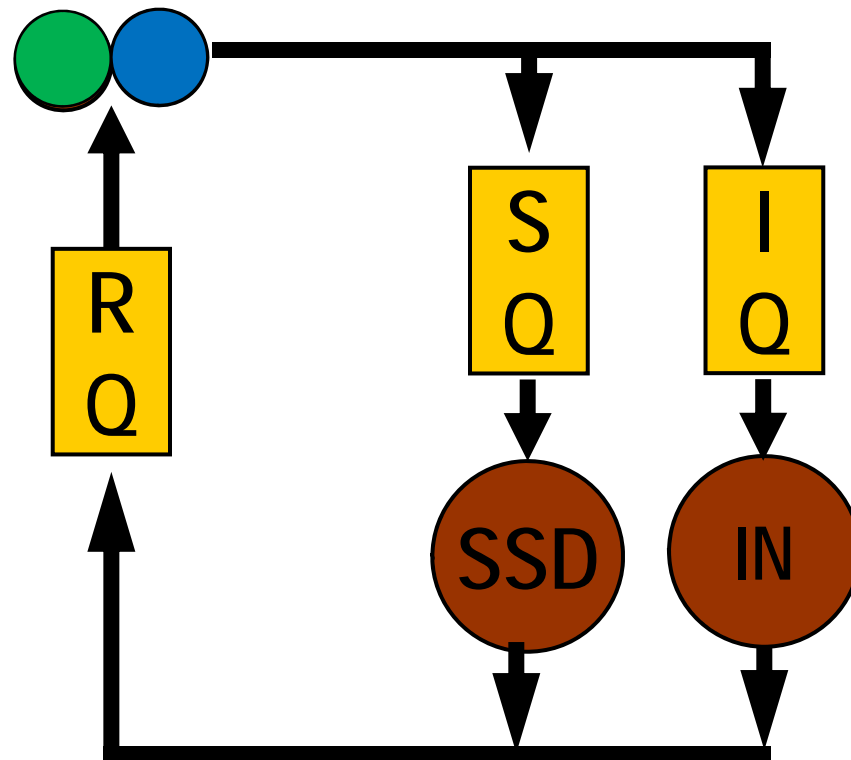
or

- Process 2 terminates at $t = 125$ ms
Process 0 is RUNNING
Process 1 is RUNNING

What happens next?

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

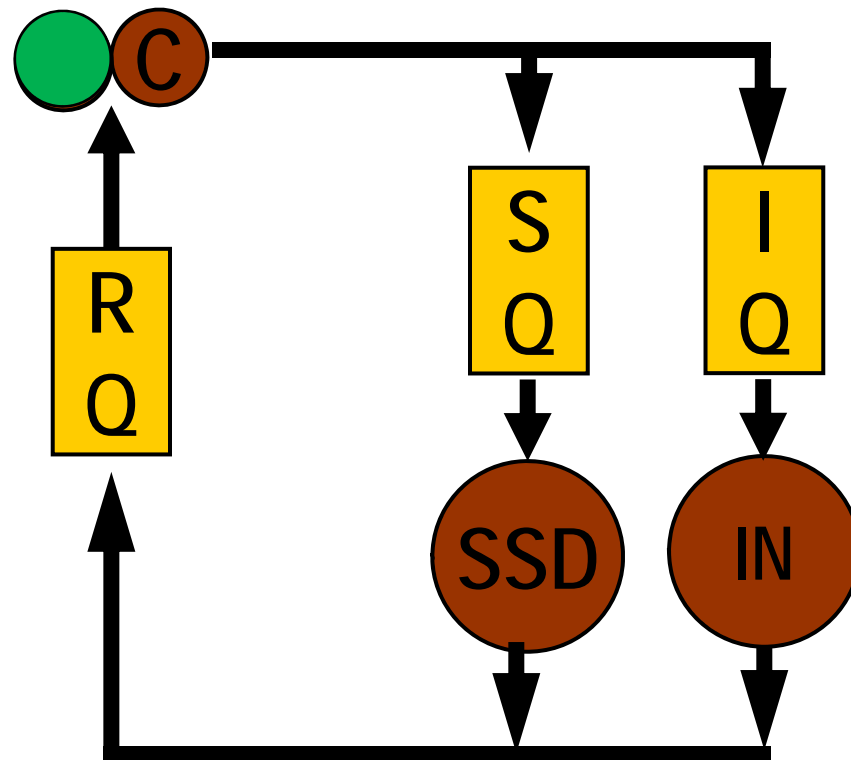
P1 holds a core until $t = 140\text{ms}$
P0 holds a core until $t = 155\text{ms}$



P1 terminates

NCORES 2
NEW 5
CORE 100
SSD 0
CORE 30
NEW 20
CORE 50
SSD 0
CORE 50
NEW 30
CORE 20
SSD 0
CORE 20

P0 holds a core until $t = 155\text{ms}$





Your program will display

- Time = 140ms
Process 1 terminates
Process 0 is RUNNING

or

- Process 1 terminates at $t = 140 \text{ ms}$
Process 0 is RUNNING

What happens next?

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50


NEW 30

CORE 20

SSD 0

CORE 20

P0 holds a core until $t = 155\text{ms}$ ←



P0 terminates

NCORES 2

NEW 5

CORE 100

SSD 0

CORE 30

NEW 20

CORE 50

SSD 0

CORE 50


NEW 30

CORE 20

SSD 0

CORE 20

P0 holds a core until $t = 155\text{ms}$







Your program will display

- Time = 155ms
Process 0 terminates

or

- Process 0 terminates at $t = 155\text{ms}$



Your program will display (II)

■ SUMMARY:

Number of processes that completed: 3

Total number of SSD accesses: 3

Average SSD access time: 0 ms

Total elapsed time: 155 ms

Core utilization: 174.19 percent

SSD utilization: 0.00 percent

$$\begin{aligned}\text{Core Utilization} &= (100 + 30 + 50 + 50 + 20 + 20)/155 \\ &= 270/155 = 1.7419\end{aligned}$$



SSD average access time (I)

- Must consider SSD contention
 - Unlikely in most cases
- Must compute difference between
 - SSD request time
 - SSD request completion timefor each SSD request



SSD average access time (II)


SSD request time

SSD
Queue

Request starts
getting processed

SSD request
completion time

SSD





The best way to implement it

- Have two counters both initialized at zero
 - **SSDcount**
 - **SSDtimes**
- When a SSD request is issued by a process
 - Add 1 to **SSDcount**
 - **Subtract** current time from **SSDtimes**
- When the SSD request is completed
 - **Add** current time to **SSDtimes**



Example

- Assume a request processing time of 5 ms
- First request arrives at time $t = 10 \text{ ms}$
- Immediately processed
- Second request arrives at time $t = 12 \text{ ms}$
- First request is completed at time $t = 15 \text{ ms}$
- Second request can be processed
- Second request is completed at time $t = 20 \text{ ms}$
- Average completion time is
 $(-10 - 12 + 15 + 20)/2 = 13/2 = 6.5 \text{ ms}$



ENGINEERING THE SIMULATION



Simulating time

- ***Absolutely nothing happens to our model between two successive "events"***
- ***Events are***
 - Arrival of a new process
 - Start of a computing step
 - Completion of a computing step
- We associate an ***event routine*** with each event



Arrival event routine

- Process first request of process
 - It will always be a CPU request



Core request routine

- **current time is clock**
request time is crt
- **if a core is free :**
mark core busy until clock + crt
else :
enter process in ready queue



Core request completion routine

- **if ready queue is empty :**
 mark core idle
 else:
 pick first process P' in ready queue
 crt' is request time for P'
 mark core busy until clock + crt'
- **proceed with next request for the process that completed the CPU request**



SSD request routine

- **current time is clock**
request time is srt
- **increment SSDcount**
- **subtract current time from SSDtimes**
- **if SSD is free :**
 mark SSD busy until clock + srt
else :
 enter process request in SSD queue



SSD request completion routine

- add current time to SSDtimes
- if SSD queue is empty :
 - mark SSD free
- else :
 - pick first process request P' in SSD queue
 - drt' is request time for P'
 - mark SSD busy until clock + srt'
- proceed with next request for process that completed its SSD request



Input request routine

- **current time is clock**
request time is irt
- **if user is free :**
 mark user busy until clock + irt
else :
 enter process request in input queue



Input request completion routine

- **if input queue is empty :**
 mark user free
- else :**
 - pick first process request P'**
 - in input queue**
 - irt' is request time for P'**
 - mark user busy until clock + irt'**
- **proceed with next request for process that completed its input request**



The simulation scheduler

1. Find next event by looking at
 - a) Core completion times
 - b) SSD termination times
 - c) Completion times of input requests
 - d) Arrival time of next process
2. Set current time to **time of next event**
3. Process event routine
4. Repeat until all processes are done



Organizing our program (I)

- Most steps of simulation involve scheduling future completion events
- Associate with each completion event an event notice
 - **Time of event**
 - **Device**
 - **Process sequence number**

177, "Core", 12



Organizing our program (II)

- Do the same with process starts
 - **Time of event**
 - **Process start**
 - **Process sequence number**

120, "Arrival", 19



Organizing our program (III)

- Process all event notices in chronological order


Release
SSD
247

Release
core
250

New
process
245

New
process
270

New
process
310



First notice to
be processed



Organizing our program (IV)

- Keep the event list sorted (*priority queue*)


New
process
245

Release
SSD
247

Release
core
250

New
process
270

New
process
310



First notice to be processed is
head of the list



Organizing our program (V)

- Overall organization of main program

**schedule first event # will be a process
start**

**while event list is not empty :
 process next event in list
print simulation results**



Organizing our event list

- ***Priority queue***
- Two kinds of entries
 - ***Computational steps completion times:***
 - Created and inserted "on the fly"
 - ***Process arrivals:***
 - Created during input phase
 - Already sorted



AN IMPLEMENTATION

- My main data structures would include:
 - Data table
 - Process table



The data table

- Stores the input data
- Line indices are used in process table

Operation	Parameter
NEW	5
CPU	10
INPUT	0
CPU	20
NEW	20
CPU	50
...	...



The process table (I)


Start Time	First Line	Last Line	Current Line	State
5	0	3	varies	varies
20	4	
...	



The process table (II)

- One line per process
 - ***Line index is process sequence number!***
- First column has start time of process
- First line, last line and current line respectively identify first line, last line and current line of the process in the input table
- Last column is for the current state of the process

A full list implementation of the data table



Very elegant but
harder to debug



Reading your input

- You must use I/O redirection
 - `assign1 < input_file`
- Advantages
 - Very flexible
 - Programmers write their code as if it was read from standard input
 - No need to mess with `fopen()`, `argc` and `argv`